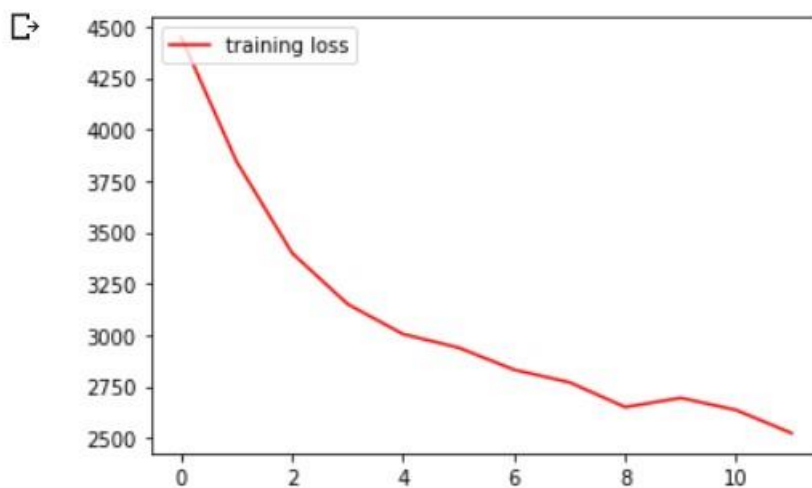


# SIT789 - Applications of Computer Vision and Speech Processing

## Credit Task 5.1: Deep learning for Computer Vision

---

```
[44] plt.plot(loss_history, label = 'training loss', color = 'r')  
plt.legend(loc = "upper left")  
plt.show()
```



```
[45] PATH = './cifar_net.pth'  
torch.save(net.state_dict(), PATH)
```

```
[46] dataiter = iter(testloader)  
images, labels = dataiter.next()  
# print images  
imshow(torchvision.utils.make_grid(images))  
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



# Time on local machine (with already using GPU on performance[gaming] laptop)

```
In [18]: import time
start_time = time.time()
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, groundtruth_labels = data
    outputs = net(images)
    _, predicted_labels = torch.max(outputs.data, 1)
    total += groundtruth_labels.size(0)
    correct += (predicted_labels == groundtruth_labels).sum().item()
print('Accuracy of the network on the 10000 test images: %d %%' % (100 * correct / total))
print("Testing time GPU is in %s seconds ---" % (time.time() - start_time))
```

Accuracy of the network on the 10000 test images: 50 %  
Testing time GPU is in 3.003812789916992 seconds ---

```
In [19]: import time
class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
start_time = time.time()
with torch.no_grad():
    for data in testloader:
        images, groundtruth_labels = data
        outputs = net(images)
        _, predicted_labels = torch.max(outputs, 1)
        c = (predicted_labels == groundtruth_labels).squeeze()
        for i in range(4):
            label = groundtruth_labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1
for i in range(10):
    print('Accuracy of %5s : %2d %%' % (classes[i], 100 * class_correct[i] / class_total[i]))
print("Testing time is in %s seconds ---" % (time.time() - start_time))
```

Accuracy of plane : 63 %  
Accuracy of car : 64 %  
Accuracy of bird : 49 %  
Accuracy of cat : 38 %  
Accuracy of deer : 43 %  
Accuracy of dog : 45 %  
Accuracy of frog : 70 %  
Accuracy of horse : 55 %  
Accuracy of ship : 44 %  
Accuracy of truck : 53 %  
Testing time is in 5.318864107131958 seconds ---

```
In [25]: import time
start_time = time.time()
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, groundtruth_labels = data
    outputs = net(images)
    _, predicted_labels = torch.max(outputs.data, 1)
    total += groundtruth_labels.size(0)
    correct += (predicted_labels == groundtruth_labels).sum().item()
print('Accuracy of the network on the 10000 test images: %d %%' % (100 * correct / total))
print("Testing time CPU is in %s seconds ---" % (time.time() - start_time))
```

Accuracy of the network on the 10000 test images: 50 %  
Testing time CPU is in 2.874325752258301 seconds ---

```
In [26]: import time
class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
start_time = time.time()
with torch.no_grad():
    for data in testloader:
        images, groundtruth_labels = data
        outputs = net(images)
        _, predicted_labels = torch.max(outputs, 1)
        c = (predicted_labels == groundtruth_labels).squeeze()
        for i in range(4):
            label = groundtruth_labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1
for i in range(10):
    print('Accuracy of %5s : %2d %%' % (classes[i], 100 * class_correct[i] / class_total[i]))

print("Testing time is in %s seconds ---" % (time.time() - start_time))
```

Accuracy of plane : 63 %  
Accuracy of car : 64 %  
Accuracy of bird : 49 %  
Accuracy of cat : 38 %  
Accuracy of deer : 43 %  
Accuracy of dog : 45 %  
Accuracy of frog : 70 %  
Accuracy of horse : 55 %  
Accuracy of ship : 44 %  
Accuracy of truck : 53 %  
Testing time is in 4.936636209487915 seconds ---

# Time on Google Colab



```
t_time = time.time()

correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, groundtruth_labels = data[0].to(device), data[1].to(device)
        outputs = net(images)
        _, predicted_labels = torch.max(outputs.data, 1)
        total += groundtruth_labels.size(0)
        correct += (predicted_labels == groundtruth_labels).sum().item()
print(device)
print('Accuracy of the network on the 10000 test images: %d %%' % (100 * correct / total))
print("Testing time is in %s seconds ---" % (time.time() - start_time))
```



cpu  
Accuracy of the network on the 10000 test images: 54 %  
Testing time is in 14.577147245407104 seconds ---



```
start_time = time.time()

correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, groundtruth_labels = data[0].to(device), data[1].to(device)
        outputs = net(images)
        _, predicted_labels = torch.max(outputs.data, 1)
        total += groundtruth_labels.size(0)
        correct += (predicted_labels == groundtruth_labels).sum().item()
print(device)
print('Accuracy of the network on the 10000 test images: %d %%' % (100 * correct / total))
print("Testing time is in %s seconds ---" % (time.time() - start_time))
```



cuda:0  
Accuracy of the network on the 10000 test images: 55 %  
Testing time is in 6.533771514892578 seconds ---

## Confusion matrix:

```
In [21]: cm = confusion_matrix(groundtruth_labels, predicted_labels)
```

```
In [22]: print(cm)
```

```
[[0 0 1 0 0 0]
 [0 0 0 0 1 0]
 [0 0 0 0 0 0]
 [0 0 0 1 0 0]
 [0 0 0 0 0 0]
 [0 0 0 0 0 1]]
```